



COMPUTATIONAL FINANCE & RISK MANAGEMENT

UNIVERSITY *of* WASHINGTON

Department of Applied Mathematics

A Probability and Statistics Primer for Quantitative Finance

Week 2: R Usage

Jake Price

Instructor, Computational Finance and Risk Management

University of Washington

Slides originally produced by Kjell Konis

Outline

Installing R and R Studio

Basic R Usage

R is an Array Programming Language

Subsetting Vectors

Installing and Loading R Packages

The R Help System

Plotting

Additional Resources

Outline

Installing R and R Studio

Basic R Usage

R is an Array Programming Language

Subsetting Vectors

Installing and Loading R Packages

The R Help System

Plotting

Additional Resources

What is R?

R is a language and environment for statistical computing and graphics

R provides . . .

- ▶ a wide variety of statistical (linear and nonlinear modeling, classical statistical tests, time-series analysis, classification, clustering, . . .) techniques
- ▶ a wide variety of graphical techniques
- ▶ and is highly extensible

R is available as Free Software under the terms of the Free Software Foundation's GNU General Public License (either Version 2 or Version 3)

R runs on Linux, Windows, and Mac OS

Installing R

The R installer is available from the *R Project* website

`http://www.r-project.org`

After installing R, run the command

```
> demo("graphics")
```

in the R Console as a quick check that things are working

What is R Studio?

R Studio is an integrated development environment (IDE) for R

R Studio includes ...

- ▶ a console
- ▶ a syntax-highlighting editor that supports direct code execution
- ▶ tools for plotting, history, debugging and workspace management

R Studio is available in an open source (i.e., free) edition – license = AGPL v3

R Studio also runs on Linux, Windows, and Mac OS

Installing R Studio

The R Studio installer is available from the R Studio website

`http://www.rstudio.com`

After installing R Studio, run the command

```
> demo("graphics")
```

in the R Studio Console as a quick check that things are working

Outline

Installing R and R Studio

Basic R Usage

R is an Array Programming Language

Subsetting Vectors

Installing and Loading R Packages

The R Help System

Plotting

Additional Resources

The R Console

The `>` symbol is called the prompt: R is prompting you to enter a command

R as a calculator: input the command `1 + 1` at the prompt and press return

```
> 1 + 1
```

The result appears on the next line

```
[1] 2
```

R can replace your 4 function calculator (R follows the usual order of operations)

+	addition	-	subtraction	^	power
*	multiplication	/	division		

Comments: any input following a `#` is ignored (until the next new line)

Mathematical Functions

R includes a wide variety of mathematical functions

For example, the `log` function returns the *natural* logarithm of its argument

```
> log(5)
```

```
[1] 1.609438
```

Some commonly used functions are given in the following table

<code>log</code>	logarithm	<code>exp</code>	exponential
<code>sin</code>	sine	<code>asin</code>	inverse sine
<code>cos</code>	cosine	<code>acos</code>	inverse cosine
<code>tan</code>	tangent	<code>atan</code>	inverse tangent
<code>sqrt</code>	square root		

Note: trigonometric functions use radians

Constants and Special Values

Constants

Exponential base: e^1

```
> exp(1)
```

```
[1] 2.718282
```

π

```
> pi
```

```
[1] 3.141593
```

Special Values

Plus and minus infinity

```
> 1/0
```

```
[1] Inf
```

```
> log(0)
```

```
[1] -Inf
```

Not a Number

```
> 0/0
```

```
[1] NaN
```

Caveats

Numerical calculations in R are done using *double precision* arithmetic

Double precision arithmetic is *not* real arithmetic

- ▶ There are a finite number of double precision numbers
- ▶ Input numbers are *coerced* to the most appropriate double precision number

Leads to *unexpected* behavior

```
> (1/3)^2 - 1 / (3^2)
```

```
[1] 0
```

But ...

```
> (1/3)^3 - 1 / (3^3)
```

```
[1] -6.938894e-18
```

Variables

Values (e.g., the output of a function) can be saved in memory, this chunk of memory is called a *variable*

For example, suppose we want to evaluate $\sin(\frac{\pi}{2})$

First, evaluate $\frac{\pi}{2}$ and store the result

```
> u <- pi / 2
```

The assignment operator in R is `<-`, it assigns the value on the right to the variable on the left

Second, compute the value of $\sin(u)$

```
> (v <- sin(u))
```

```
[1] 1
```

Variable Names

Variable names need/should conform to the following guidelines

- ▶ Variable names need to start with a character
- ▶ Variable names can contain characters, numbers, periods, and underscores
- ▶ Variable names cannot be an R reserved word

if	else	repeat	while	function
for	in	next	break	TRUE
FALSE	NULL	Inf	NaN	NA
...				

- ▶ Best practice: variable names should not be the name of a common R function
 - ▶ Some common collisions are `c`, `t`, `mean`, `var`, and `cov`

Workspace Management

Variables are stored in R's workspace

Use the `ls` function to list the variables in the workspace

```
> ls()
```

```
[1] "u" "v"
```

Use the `rm` function to remove variables from the workspace

```
> rm(u)
```

```
> ls()
```

```
[1] "v"
```

R asks to save the workspace when quitting, saved workspaces are automatically restored the next time R starts

Outline

Installing R and R Studio

Basic R Usage

R is an Array Programming Language

Subsetting Vectors

Installing and Loading R Packages

The R Help System

Plotting

Additional Resources

What is an Array Programming Language?

An array programming language generalizes operations to apply transparently to single and multi-dimensional arrays

In R ...

- ▶ a one dimensional array is called a *vector*
- ▶ a two dimensional array is called a *matrix*
- ▶ Higher dimensional arrays are possible but we are not going to use them

A goal of array programming is to provide a syntax similar to mathematical notation

- ▶ Very useful for *programming with data*

Vectors

A *vector* is an ordered collection of n values with the same underlying type

Example: suppose the vector $\mathbf{u} = (u_1, u_2)$ describes a point in the xy -plane

- ▶ Both u_1 and u_2 should be real numbers
- ▶ Since the first value defines the x coordinate and the second value the y coordinate, the vector (u_1, u_2) is (in general) different than the vector (u_2, u_1)

The number of values in the vector is called the *length* of the vector

Supported types include

- ▶ numeric
- ▶ character
- ▶ ...

Creating Vectors

The function `c` combines values into a vector

```
> (primes <- c(2, 3, 5, 7))
```

```
[1] 2 3 5 7
```

The function `length` returns the length of the input vector

```
> length(primes)
```

```
[1] 4
```

The function `seq` creates sequences of equally-spaced values

```
> seq(-1, 1, 0.5)
```

```
[1] -1.0 -0.5  0.0  0.5  1.0
```

Naming Vector Components

The = operator assigns a name to a value

Names can be given to the components of a vector

```
> wall.street.wiz <- c(houses = 3, cars = 5, boats = 2)
```

```
> wall.street.wiz
```

houses	cars	boats
3	5	2

The names function returns the names of the components of a vector

```
> names(wall.street.wiz)
```

```
[1] "houses" "cars"   "boats"
```

Remark: assigning a name to a value (=) often has the same effect as assigning a value to a variable (<-) \implies = and <- can be used interchangeably

Matrices

A *matrix* is a two dimensional array of values (of the same type) described by a number of rows and a number of columns

The function `cbind` combines vectors (with the same length) into a matrix

```
> col1 <- c(11, 21, 31, 41)
> col2 <- c(12, 22, 32, 42)
> (mat <- cbind(col1, col2))
```

	col1	col2
[1,]	11	12
[2,]	21	22
[3,]	31	32
[4,]	41	42

Component-wise Operations on Vectors

What does it mean to

... generalize operations to apply transparently to single and multi-dimensional arrays

When a vector is provided as the argument of a function, the output is a vector containing the function applied to each component of the input vector

```
> (x <- 1:9) # x <- seq(1, 9, 1)
```

```
[1] 1 2 3 4 5 6 7 8 9
```

```
> sqrt(x)
```

```
[1] 1.000 1.414 1.732 2.000 2.236 2.449 2.646 2.828 3.000
```

Component-wise Composite Functions

Shifting: adding a length n vector and a length 1 vector

```
> x
```

```
[1] 1 2 3 4 5 6 7 8 9
```

```
> x + 2
```

```
[1] 3 4 5 6 7 8 9 10 11
```

Scaling: multiplying a length n vector by a length 1 vector

```
> 3*x
```

```
[1] 3 6 9 12 15 18 21 24 27
```

Element-wise Composite Functions

Functions act element-wise on matrices

```
> mat
```

	col1	col2
[1,]	11	12
[2,]	21	22
[3,]	31	32
[4,]	41	42

```
> 3*mat + 2
```

	col1	col2
[1,]	35	38
[2,]	65	68
[3,]	95	98
[4,]	125	128

Outline

Installing R and R Studio

Basic R Usage

R is an Array Programming Language

Subsetting Vectors

Installing and Loading R Packages

The R Help System

Plotting

Additional Resources

Logical-Valued Functions

Logical values in R are TRUE and FALSE

Functions with a yes/no output generally return logical values

```
> x
```

```
[1] 1 2 3 4 5 6 7 8 9
```

```
> x > 4
```

```
[1] FALSE FALSE FALSE FALSE  TRUE  TRUE  TRUE  TRUE  TRUE
```

```
> (x.gt.eq.4 <- x >= 4)
```

```
[1] FALSE FALSE FALSE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
```

```
> x == 4  # Remember the double precision caveat
```

```
[1] FALSE FALSE FALSE  TRUE FALSE FALSE FALSE FALSE FALSE
```

Logical Vectors

The variable `x.gt.eq.4` is a logical vector, that is

- ▶ it is a vector
- ▶ the underlying type of value is logical

Lets create a second logical vector

```
> (x.lt.eq.7 <- x <= 7)
```

```
[1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE FALSE FALSE
```

Operations on Logical Vectors

Three key operations for logical vectors: & (AND), | (OR), and ! (NOT)

Suppose that x and y are logical vectors with the same length

AND: `ans <- x & y`

`ansi` is TRUE if both `xi` and `yi` are TRUE, otherwise `ansi` is FALSE

```
> x.gt.eq.4
```

```
[1] FALSE FALSE FALSE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
```

```
> x.lt.eq.7
```

```
[1]  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE FALSE FALSE
```

```
> x.gt.eq.4 & x.lt.eq.7
```

```
[1] FALSE FALSE FALSE  TRUE  TRUE  TRUE  TRUE FALSE FALSE
```

Operations on Logical Vectors (continued)

OR: `ans <- x | y`

`ansi` is FALSE if both `xi` and `yi` are FALSE, otherwise `ansi` is TRUE

```
> x.gt.eq.4
```

```
[1] FALSE FALSE FALSE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
```

```
> x.lt.eq.7
```

```
[1]  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE FALSE FALSE
```

```
> x.gt.eq.4 | x.lt.eq.7
```

```
[1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
```

Operations on Logical Vectors (continued)

NOT: `ans <- !x`

`ansi` is FALSE if `xi` is TRUE, otherwise `ansi` is TRUE

```
> x.gt.eq.4
```

```
[1] FALSE FALSE FALSE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
```

```
> !x.gt.eq.4
```

```
[1]  TRUE  TRUE  TRUE FALSE FALSE FALSE FALSE FALSE FALSE
```

Subsetting Vectors

The subset operator in R is square brackets: `[]`

If

- ▶ `x` is a vector (with any underlying type)
- ▶ `idx` is a logical vector the same length as `x`

then `x[idx]` is a vector (with the same underlying type as `x`) containing the components of `x` for which `idx` is `TRUE`

```
> x[x.gt.eq.4]
```

```
[1] 4 5 6 7 8 9
```

```
> x[x.lt.eq.7]
```

```
[1] 1 2 3 4 5 6 7
```

Subsetting Vectors – More Examples

```
> x[x < 6]
```

```
[1] 1 2 3 4 5
```

```
> x[x > 2 & x < 8]
```

```
[1] 3 4 5 6 7
```

```
> x[x <= 2 | x >= 8]
```

```
[1] 1 2 8 9
```

```
> x[!(x == 3 | x == 6)]
```

```
[1] 1 2 4 5 7 8 9
```


Subsetting Vectors by Index

Suppose that x is a vector of length n

Let idx be a vector of values from the set $\{1, 2, \dots, n\}$

The vector $ans \leftarrow x[idx]$ is the same length as idx and has components

$$ans_i = x_{idx_i}$$

```
> primes
```

```
[1] 2 3 5 7
```

```
> primes[c(2, 4)]
```

```
[1] 3 7
```

From now on: write $ans[i]$ to mean ans_i

Outline

Installing R and R Studio

Basic R Usage

R is an Array Programming Language

Subsetting Vectors

Installing and Loading R Packages

The R Help System

Plotting

Additional Resources

R Packages

An R package is a collection of functions and data (and ...)

⇒ loading a package into your R session extends the capabilities of R

Packages loaded by default

```
> search()
```

```
[1] ".GlobalEnv"          "package:stats"      "package:graphics"
[4] "package:grDevices"  "package:utils"      "package:datasets"
[7] "package:methods"    "Autoloads"          "package:base"
```

For the curious: use the `find` function to see what package provides an object

```
> find("log")
```

```
[1] "package:base"
```

Loading R Packages

One useful function is the `fitdistr` function from the MASS package

MASS is a recommended package, it is already installed on your computer

The `library` function loads and attaches packages to the R session

```
> library(MASS)
```

The functions and data provided by the MASS package are now available to the R session

```
> search()
```

```
[1] ".GlobalEnv"          "package:MASS"        "package:stats"
[4] "package:graphics"    "package:grDevices"   "package:utils"
[7] "package:datasets"    "package:methods"     "Autoloads"
[10] "package:base"
```

```
> find("fitdistr")
```

```
[1] "package:MASS"
```

Installing R Packages

The extensive collection of user-contributed packages is one of the reasons for the popularity of R

The Comprehensive R Archive Network (CRAN) is a clearing house for user-contributed packages

The `install.packages` function downloads packages from CRAN and installs them on your computer

The `getSymbols` function on the `quantmod` package will be the primary source of price data used in this course

```
> install.packages("quantmod")
```

The package must be loaded and attached to the R session

```
> library("quantmod")
```

```
> find("getSymbols")
```

```
[1] "package:quantmod"
```

Summary R Packages

Installing a package means downloading the package from CRAN and installing it on the computer

Installing only needs to be done once (per computer)

Loading and attaching a package means telling the R session to look for functions (data, etc.) in the package

Loading and attaching needs to be done in each R session

Outline

Installing R and R Studio

Basic R Usage

R is an Array Programming Language

Subsetting Vectors

Installing and Loading R Packages

The R Help System

Plotting

Additional Resources

Getting Help

The `log` function returns the natural logarithm of its argument

How were we supposed to know that?

The `help` function displays the documentation associated with an object

```
> help(log)
```

Documentation is divided into sections including

Description a description of the function or functions

Usage how to call the function

Arguments supported types of arguments

Details additional explanation if needed

Value a description of the returned object

See Also related functions

Examples

Description

``log'` computes logarithms, by default natural logarithms, ``log10'` computes common (i.e., base 10) logarithms, and ``log2'` computes binary (i.e., base 2) logarithms. The general form ``log(x, base)'` computes logarithms with base ``base'`.

``log1p(x)'` computes $\log(1+x)$ accurately also for $|x| \ll 1$.

``exp'` computes the exponential function.

``expm1(x)'` computes $\exp(x) - 1$ accurately also for $|x| \ll 1$.

Usage

```
log(x, base = exp(1))
```

```
logb(x, base = exp(1))
```

```
log10(x)
```

```
log2(x)
```

```
log1p(x)
```

```
exp(x)
```

```
expm1(x)
```

Arguments

`x`: a numeric or complex vector.

`base`: a positive or complex number: the base with respect to which logarithms are computed. Defaults to `e=exp(1)`.

Details

All except ``logb'` are generic functions: methods can be defined for them individually or via the ``Math'` group generic.

``log10'` and ``log2'` are only convenience wrappers, but logs to bases 10 and 2 (whether computed `_via_`log'` or the wrappers) will be computed more efficiently and accurately where supported by the OS. Methods can be set for them individually (and otherwise methods for ``log'` will be used).

``logb'` is a wrapper for ``log'` for compatibility with S. If (S3 or S4) methods are set for ``log'` they will be dispatched. Do not set S4 methods on ``logb'` itself.

All except ``log'` are primitive functions.

Value

A vector of the same length as ``x'` containing the transformed values. ``log(0)'` gives ``-Inf'`, and ``log(x)'` for negative values of ``x'` is ``NaN'`. ``exp(-Inf)'` is ``0'`.

For complex inputs to the log functions, the value is a complex number with imaginary part in the range $[-\pi, \pi]$: which end of the range is used might be platform-specific.

Outline

Installing R and R Studio

Basic R Usage

R is an Array Programming Language

Subsetting Vectors

Installing and Loading R Packages

The R Help System

Plotting

Additional Resources

The plot Function

The `plot` function produces scatter plots

Usage:

```
plot(x, y, ...)
```

The points are (x_i, y_i) , the argument `x` is a vector containing the `x` values and `y` is a vector (the same length as `x`) containing the `y` values


The `...` is a placeholder for optional arguments

Example:

```
> x <- seq(0.1, 5, 0.1)
> y <- log(x)
```


Plot with Defaults

```
> plot(x, y)
```



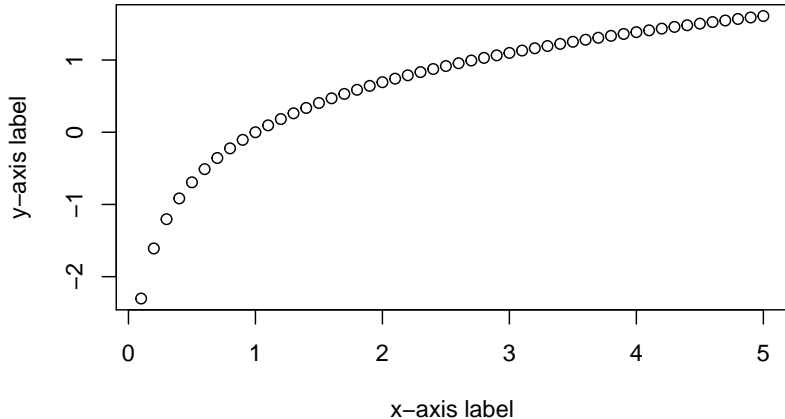
Optional Arguments: type

```
> plot(x, y, type = "l")
```



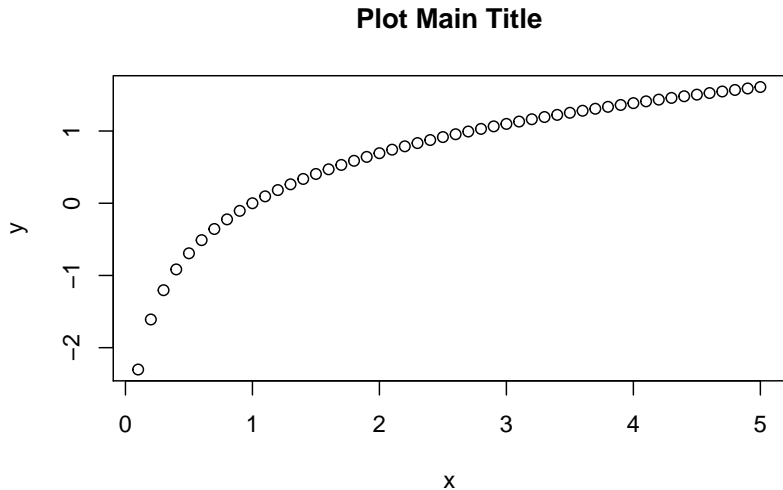
Optional Arguments: Axis Labels

```
> plot(x, y, xlab = "x-axis label", ylab = "y-axis label")
```




Optional Arguments: Title

```
> plot(x, y, main = "Plot Main Title")
```




Optional Arguments: Plotting Region

```
> plot(x, y, xlim = c(-1, 6), ylim = c(-3, 3))
```




Optional Arguments: Setting the Aspect Ratio

```
> plot(x, y, asp = 1)
```




Optional Arguments: Plotting Symbol

```
> plot(x, y, pch = 16)
```




Optional Arguments: Color

```
> plot(x, y, col = "blue")
```



Square Plotting Region

```
> par(pty = "s")  
> plot(x, y)
```



Object-Oriented Plotting

R is an object-oriented programming language \implies behavior of the plot function depends on what `x` is

If `x` is a vector then we get a scatter plot

If `x` is a financial prices series


```
> getSymbols("SBUX")
```

```
[1] "SBUX"
```

The variable `SBUX` contains 10 years of price data for Starbucks

Plotting a Price Series

```
> plot(SBUX)
```



Outline

Installing R and R Studio

Basic R Usage

R is an Array Programming Language

Subsetting Vectors

Installing and Loading R Packages

The R Help System

Plotting

Additional Resources

Additional Resources for Learning R

DataCamp: The easiest way to learn R programming and data science

<https://www.datacamp.com>



COMPUTATIONAL FINANCE & RISK MANAGEMENT

UNIVERSITY *of* WASHINGTON

Department of Applied Mathematics

<http://computational-finance.uw.edu>